

CSE 2421

Integer Multiplication and Division

Base 10 Multiplication Review (obviously)

a (multiplicand) * b (multiplier)

Consider a = 123, b = 123.

$$\begin{array}{r} 123 \\ \times 123 \\ \hline 369 \text{ (this is } 123 * 3) \\ 246 \text{ (This is } 123 * 2 \text{ shifted left 1 place)} \\ + 123 \text{ (This is } 123 * 1 \text{ shifted left 2 places)} \\ \hline 15129 \end{array}$$

Binary (Base 2) Multiplication

Now consider:

a (multiplicand) * b (multiplier) where a = 1011 (11 decimal), b = 1110 (14 decimal)

				1	0	1	1		(this is 11 in decimal)			
			x	1	1	1	0		(this is 14 in decimal)			
			=====									
				0	0	0	0		(this is 1011 x 0)			
				1	0	1	1		(this is 1011 x 1, shifted one position to the left)			
			1	0	1	1			(this is 1011 x 1, shifted two positions to the left)			
+	1	0	1	1					(this is 1011 x 1, shifted three positions to the left)			
			=====									
	1	0	0	1	1	0	1	0	(this is 154 in decimal)			

This is the same process that we use for decimal multiplication!

One could say that binary multiplication is actually easier because each row is either just a “shifted copy” of the multiplicand or it’s 0.

Hardware for integer multiplication (Simplest version)

- Suppose we want the CPU to multiply two operands:
 a (multiplicand) * b (multiplier)
- Typically, in the simplest version of multiplication hardware, 4 additional registers are used:
 - Multiplicand register (contains a copy of the multiplicand)
 - Multiplier register (contains a copy of the multiplier)
 - Shifted multiplicand register
 - Result register

Multiplication hardware cont

- For n bit multiplication, we can denote the bits in the multiplier as b_0 (least significant bit) to b_{n-1} (most significant bit)
- The hardware initializes the result register to 0
- Then, the hardware does the following (pseudo-code):

```
for (j = 0; j < n; j++) {  
    if ( $b_j == 1$ ) copy multiplicand register contents, left shifted by  $j$   
    bits, to shifted multiplicand register;  
    else copy 0 to shifted multiplicand register;  
    Add shifted multiplicand register to result register;  
}
```

Integer multiplication

- B2U 8-bit range → 0 to 255

	8-bit multiplication												
12 ->					0	0	0	0	1	1	0	0	
9 ->					0	0	0	0	1	0	0	1	
					<hr/>								
					0	0	0	0	1	1	0	0	
	0	0	0		0	1	1	0	0				
													<- carry
					<hr/>								
					0	1	1	0	1	1	0	0	= 108

Integer multiplication

- B2T 8-bit range \rightarrow -128 to 127

	8-bit multiplication													
-4 ->					1	1	1	1	1	1	0	0		
9 ->					0	0	0	0	1	0	0	1		
					1	1	1	1	1	1	0	0		
	1 1 1				1	1	1	0	0					
	1 1 1 1				1	1							<- carry	
	1	0	0	0	1	1	0	1	1	1	0	0	= -36	

Integer multiplication (cont)

- Unsigned i.e. simple binary

- For x and y , each with the same width (w)
- $x*y$ yields a w -bit value given by the w bits of the integer product
 - ✓ Result interpreted as an unsigned value
 - ✓ Overflow occurs if the result will not fit in w bits; to store any possible result for multiplication of two w bit values, $2w$ bits are needed.

- Signed = similar, but result interpreted as signed value

Overflow can occur under same conditions as above.

Integer multiplication (4 bits)

binary	unsigned	two's comp	result
0111*0011	7*3=21=0001 0101	same	0101
	21 mod 16 = 5	same	
1001*0100	9*4=36=0010 0100	-7*4=-28=1110 0100	0100
fyi	36 mod 16 = 4	-28 mod 16 = 4	
1100*0101	12*5=60=0011 1100	-4*5=-20=1110 1100	1100
	60 mod 16 = 12	-20 mod 16 = -4	
1101*1110	13*14=182=1011 0110	-3*-2=6=0000 0110	0110
	182 mod 16 = 6	6 mod 16 = 6	
1111*0001	15*1=15=0000 1111	-1*1=-1=1111 1111	1111
	15 mod 16 = 15	-1 mod 16 = -1	

The world has changed

- ▶ The imul instruction on current intel architectures takes 1 clock cycle and has latency of about 3 clock cycles.
- ▶ A shift is 1 clock cycle with a latency of 1 clock cycle
- ▶ The transistor budget for modern CPUs are *huge* compared to previous generations of chips. It takes a *lot* of adders to make a large, fast multiplier
- ▶ The intel 8086 took 100 clock cycles for imul
- ▶ Other chips may have high costs to multiply compared to the cost of shifting
- ▶ ARM 9 chips need 2–4 clock cycles to multiply

Multiply by constants (known at compile time)

- First case: Multiplying by a power of 2
 - Power of 2 represented by k
 - This is an optimization for the C compiler: left shifting can be used to replace multiplication, which is *used to be a much* slower operation on small chips.
(8 bit Arduino does 8 bit multiply in 1-2 clocks, 16 bit in 6 clocks, 32 bit in 12 clocks)
 - So k zeroes added in to the right side of x
 - Shift left by k : $x \ll k$
 - Overflow issues the same as $x * y$
- What is $x * 4$ where $x = 5$?
 - $x = 5 = 00000101$
 - $4 = 2^k$, and $k = 2$
 - $x \ll k = 00010100 = 20$

Multiply by constants (known at compile time)

- First case: Multiplying by a power of 2
 - Power of 2 represented by k
 - This is an optimization for the C compiler: left shifting can be used to replace multiplication, which is *used to be a much* slower operation on small chips.
(8 bit Arduino does 8 bit multiply in 1-2 clocks, 16 bit in 6 clocks, 32 bit in 12 clocks)
 - So k zeroes added in to the right side of x
 - Shift left by k : $x \ll k$
 - Overflow issues the same as $x * y$
- What is $x * 4$ where $x = 5$?
 - $x = 5 = 00000101$
 - $4 = 2^k$, and $k = 2$
 - $x \ll k = 00010100 = 20$
- What if $x = -5 = 1111\ 1011$
 - We can shift a maximum of 4 times before we lose the sign bit

Multiply by constants

- General case
 - Every binary value is an addition of powers of 2
 - Shifts, adds, and subtracts generally take 1 clock cycle, **many** fewer clock cycles than multiplication on chips with a constrained transistor count (ARM takes 2-4 cycles)
 - This is an optimization for the C compiler
 - For the compiler to do this kind of optimization, one operand must be a constant known at compilation time
 - Has to be a run of consecutive 1's to work (to save significant CPU cycles)
 - Where n = position of leftmost 1 bit in the run and m = the rightmost (with bit positions designated as $w-1$ for msb to 0 for lsb)
- Example: suppose for constant k , $k = 7 = 0111$ /* $n=2, m=0$ */
- Using the 1's in the constant (with addition): $2^2 + 2^1 + 2^0 = 7$
 - $(x \ll n) + (x \ll n-1) + \dots + (x \ll m)$ /* use the 1's */
- Using the 0 in the constant (with subtraction) : $2^3 - 2^0 = 8 - 1 = 7$
 - $(x \ll n+1) - (x \ll m)$ /* use the 0 */

Multiply by constants

```
x = 5;          /* 0101 */
x*7 = 35        /* 7=0111: n = 2 and m = 0 */
x<<2 + x<<1 + x<<0
   00010100
+ 00001010
+ 00000101
   00100011 = 35
```

Multiply by constants

```
x = 5;          /* 0101 */  
x*7 = 35        /* 7=0111: n = 2 and m = 0 */
```

```
x<<2 + x<<1 + x<<0
```

```
  00010100  
+ 00001010  
+ 00000101  
  00100011 = 35
```

OR

```
x<<3 - x<<0  
  00101000  
- 00000101  
  00100011 = 35
```

Multiply by constants (cont)

- What if the bit position n in the constant is the most significant bit?
- Since the formula with subtraction is $(x \ll n+1) - (x \ll m)$ and shifting $n+1$ times gives zero, then the formula gives $-(x \ll m)$
- Example: Assuming 4 bit values 2^3 is the most significant bit (that is, the msb position is 3). If we were to shift something 4 times it effectively 0's all 4 bits

Multiply by constants (cont)

- Assuming 4 bit values

```
x = 5; /*0101*/
```

```
y=12; /*1100 so n=3, m=2 */
```

```
z = x*y;
```

```
x<<4 - x<<2
```

```
0000          /*8 bits: 01010000 */
```

```
- 0100        /*8 bits: 00010100 */
```

```
1100 = -4    /*8 bits: 00111100 - CORRECT!*/
```

Overflowed, because information in high-order bits was lost!

We will look more at overflow due to shifting below.

Multiply by constants (cont)

- What if y is negative?

- Remember original formula (with subtraction): $(x \ll n+1) - (x \ll m)$ where n =position of leftmost and m =position of rightmost 1 in the constant

- Negate formula for negatives (multiply by -1)**, and **negate the constant** (that is, **treat the constant as a positive**):

- Formula for negatives: $(x \ll m) - (x \ll n+1)$

- Instead of subtracting, we can take the 2's complement of $x \ll n+1$, and add:

- $(x \ll m) + (- (x \ll n+1))$

- Example: $y = -6$ $x = 5$ /*Assuming 8-bits*/

- $x * y$

- $y = 0000\ 0110 = +6$ /*treat constant as positive: $n=2$ and $m=1$ */

- $2^1 - 2^3 = 2 - 8 = -6$

- $(x \ll 1) - (x \ll 3)$

- 0000 1010 ($x=5$, shifted 1 left)

- 0010 1000 ($x=5$, shifted 3 left)

- 1110 0010 = -30

Overflow for unsigned multiplication

- In CPUs, the operands and the result of the multiplication are normally stored in the same number of bits
- To eliminate the possibility of overflow for multiplication of w bit operands, we need $2w$ bits.
- If we only have w bits to store the result, can we describe when overflow will occur due to shifting? [Note that the addition operations may also cause overflow!]
- Using the same kind of notation that we used for the bits in constants above, if we designate the position of the most significant 1 in $op1$ as $op1_n$, and the position of the most significant 1 in $op2$ as $op2_n$, when will we get overflow due to shifting?
- Let's think about the algorithm we described earlier – we need to answer two questions to describe the conditions under which overflow due to shifting will occur for unsigned operands.

Overflow for unsigned multiplication

- Question 1:

- If the most significant bit in $op1$ is in position $op1_n$, and if operands and the result have w bit representations, how many bits can $op1$ be shifted left before overflow due to shifting occurs?

- Question 2:

- What is the relationship between $op2_n$ and the maximum number of bits that $op1$ is shifted left by the multiplication algorithm?

Overflow for unsigned multiplication

- Question 1:

- If the most significant bit in $op1$ is in position $op1_n$, and if operands and the result have w bit representations, how many bits can $op1$ be shifted left before overflow due to shifting occurs?

- Answer: $w - 1 - op1_n$

- Example: $w=4$, $op1_n=2$ (0100), then $4-1-2 = 1$

- we can only shift left once; any more than that causes overflow

- Question 2:

- What is the relationship between $op2_n$ and the maximum number of bits that $op1$ is shifted left by the multiplication algorithm?

- Answer: $op2_n$ is the maximum number of bits that $op1$ is shifted left

Overflow for unsigned multiplication

- Conclusion: If $op2_n > w - 1 - op1_n$, overflow due to shifting will occur

So if $op2_n > 1$ in our example, overflow occurs.

- Again, overflow may also occur due to addition, but the result above describes cases where overflow will *definitely* occur due to shifting.

Example

- ▶ $w=8, op_1=4, op_2=3$

- ▶ $w(8) - 1 - op_1(4) = 3$

Example

- ▶ $w=8, op_1=4, op_2=3$
- ▶ $w(8) - 1 - op_1(4) = 3$
- ▶ Is $op_2 > w - 1 - op_1$? $3 > 3$?

Example

- ▶ $w=8, op_1=4, op_2=3$
- ▶ $w(8) - 1 - op_1(4) = 3$
- ▶ Is $op_2 > w - 1 - op_1$? $3 > 3$?
- ▶ No overflow in this case.

Dividing by powers of 2

- Even slower than integer multiplication
- Dividing by powers of 2 → right shifting
 - There are 2 types of right shifting:
 - Logical – unsigned: fill with 0's
 - Arithmetic – two's complement: fill with copy of msb
- Integer division always truncates
 - C float-to-integer casts round towards zero.
 - Division by right shifting rounds *down*.
 - These rounding errors generally accumulate

Unsigned Integer Division

- Dividing by 2^k
- Logical right shift by k ($x \gg k$)
- Then rounding down (toward zero)
- Remember that logical shift will shift in zeroes to fill the most significant bits when the original bit pattern is shifted right

Dividing by powers of 2

8÷2	1	0	0	9÷2	1	0	0	12÷4		1	1	15÷4		1	1				
10	1	0	0	0	10	1	0	0	1	100	1	1	0	0	100	1	1	1	1
	1	0				1	0				1	0	0			1	0	0	
		0	0				0	1			1	0	0				1	1	1
		0	0				0	0			1	0	0				1	0	0
		0	0					1					0				1	1	

$$x = 8 = 1000$$

$$y = 2 = 0010 = 2^1$$

$$x \gg 1 = 100 \quad (x / y)$$

$$x = 9 = 1001$$

$$y = 2 = 0010 = 2^1$$

$$x \gg 1 = 100 \quad (x / y)$$

Dividing by powers of 2

8÷2	1	0	0		9÷2	1	0	0		12÷4		1	1		15÷4			1	1
10	1	0	0	0	10	1	0	0	1	100	1	1	0	0	100	1	1	1	1
	1	0				1	0				1	0	0			1	0	0	
			0	0			0	1				1	0	0			1	1	1
			0	0			0	0				1	0	0			1	0	0
			0	0				1					0				1	1	

$$x = 12 = 1100$$

$$y = 4 = 0100 = 2^2$$

$$x \gg 2 = 11$$

$$x = 15 = 1111$$

$$y = 4 = 0100 = 2^2$$

$$x \gg 2 = 11$$

Signed Integer Division

- Two's complement
- Sign extend for arithmetic shift (fill with copy of msb)
- Rounds down (away from zero for negative results)
- $-7/2$ will yield -4 rather than -3
- $x = -82 \quad y = 2^k$

k	Binary	Decimal	$-82/2^k$
0	1 0 1 0 1 1 1 0	-82	-82.000000
1	1 1 0 1 0 1 1 1	-41	-41.000000
2	1 1 1 0 1 0 1 1	-21	-20.500000
3	1 1 1 1 0 1 0 1	-11	-10.250000
4	1 1 1 1 1 0 1 0	-6	-5.125000
5	1 1 1 1 1 1 0 1	-3	-2.562500

Note rounding effect