

Performance

This is only the introduction

Required reading: Chapter 5:
Introduction and 5.1, plus the first paragraphs of
sections 5.2 – 5.9. Look over 5.15

C/assembly



How Fast?

Godot TBM owned by The Boring Company

Gary, The Boring Company's snail mascot

Usain Bolt

Smart For2

Tesla Model S P100DL

Jetliners

Falcon 9 / Long March to orbit

Earth Escape velocity

Photons



How Fast?

Godot TBM owned by The Boring Company (300' / week?)

Gary, The Boring Company's snail mascot (14x faster)

Usain Bolt (27 MPH / 44 km/hr)

Smart For2 (96 MPH)

Tesla Model S P100DL (155 MPH)

Jetliners [.87 to .93 Mach] (Mach 1 is 760 MPH / 1200 km/hr)


Falcon 9 / Long March to orbit (Mach 20)

Earth Escape velocity (Mach 33)

Photons at c



How Quick? (0–60mph / 0–100)

- ▶ Godot?
 - ▶ Gary the snail?
 - ▶ Usain Bolt?
 - ▶ Smart For2?
 - ▶ Tesla Model S P100DL?
 - ▶ Boeing 777?
 - ▶ Falcon 9?
 - ▶ Photon?
- 

How Quick? (0–60mph / 0–100)

- ▶ Godot – N/A
- ▶ Gary – N/A
- ▶ Usain Bolt – (N/A)
- ▶ Smart For2 – 11 sec
- ▶ Tesla Model S P100DL – 2.5 sec
- ▶ Boeing 777 (unladen) 6 sec, most jetliners can make it to 150 MPH in 30 seconds or less.
- ▶ Falcon 9 in ~ 10–11 seconds
- ▶ Photon to c in 0 seconds

What do those measures mean?

- ▶ Categories of performance
 - “too slow”
 - “OK”
 - “ridiculously fast”
- ▶ Sufficient Performance
 - Meets need
 - Doesn't meet the need


Measures

- ▶ We can't talk about performance without talking about measures
- ▶ If you don't measure, you don't know
- ▶ If you don't profile, you can't tell where

Knuth

- ▶ Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We should forget about small efficiencies, say about 97% of the time: **premature optimization is the root of all evil**. Yet we should not pass up our opportunities in that critical 3%.

How fast is determined by timing your code

- ▶ Different languages have access to different clock functions
 - 1 second resolution
 - Sub-second resolution
 - ▶ Use stats when your clock is too coarse-grained.
 - ▶ Test code should include timing data as an output.
- 

Bottlenecks are found by profiling

- ▶ Profile optimized code, not debug code
 - Do not check performance on debug code compiled with `-g`
 - Only check performance on optimized code compiled with `-O`
 - To enable profiling compile with `-pg` then use `gprof`
- ▶ `gprof` is the finest technology from 1980's – cell phones back then were called “brick phones”
- ▶ Newer tools have newer profiling capabilities that are platform dependent.

Systems need performance checking, too!

- ▶ top command under Unix
- ▶ task manager and perfmon under windows

Which is faster? Why? Does it matter?

- ▶ `Sum = 0;`
- ▶ `For (i=0; i<limit; i++) Sum = sum + a[i];`

- ▶ `Sum = 0;`
- ▶ `Ptr = a;`
- ▶ `For(i=0; i<limit;i++) Sum = sum + *ptr++;`

Which is faster? Why? Does it matter?

```
Sum = 0;
for (i=0; i<rows; i++)
{
    for (k=0; k<cols; k++)
    {
        Sum += a[i][k];
    }
}
```

C is row-major order
Fortran is column-major order

```
Sum = 0;
for (i=0; i<rows; i++)
{
    for (k=0; k<cols; k++)
    {
        Sum += a[k][i];
    }
}
```

Java uses 1D vectors instead

Performance At Scale

- ▶ AWS Whitepaper from 2014
 - <https://www.enterpriseai.news/2014/11/14/rare-peek-massive-scale-aws/?fbclid=IwAR2swFdPs5PlfAg8rS3uce95y5VRskr81TXE4kfJjOtTaIHApSPUwuZ3SL0>
- ▶ Availability zones in a region are always under 2ms apart, usually under 1 ms apart (186 miles) – the write time of an SSD device.
- ▶ Data centers in an availability zone are under ¼ ms apart (46 miles).
- ▶ A data center is over 50,000 servers in 2,000 racks using 25–30 MW of power
- ▶ Custom network stack on custom hardware

What Do Those Numbers Mean?

- ▶ New York to LA is 74ms latency on a real network
- ▶ Ground to geosynchronous orbit is 120ms one way at light speed
- ▶ Starlink is 340 miles up (1.83ms one way, 41ms actual average ping)
- ▶ High performance computer systems and networks are about latency as well as throughput, so we measure distances in terms of light-speed delays to get a minimum delay knowing that real networks take longer
- ▶ Inside a data center, software latency is far worse than light-time lag:
 - Milliseconds to get from the application to the network card
 - Microseconds to get through the card
 - Nanoseconds to transit the fiber

But does it matter? Python→C speedup

The program does matrix multiplication

- ▶ Switching to C: 47x speedup
- ▶ Parallel loops on many cores: 7x speedup
- ▶ Optimal memory layout for cache: 20x speedup
- ▶ Using SIMD floating point extensions: 7x speedup

Total speedup: 62,000 faster!

Un-asked questions:

- Did anyone care?
- Was it worth the effort?